

03.3-Regression Lineal PyTorch

November 18, 2019

```
In [1]: from pylab import *

In [2]: from sklearn.datasets import load_boston

In [3]: boston_dataset = load_boston()

In [4]: import pandas as pd

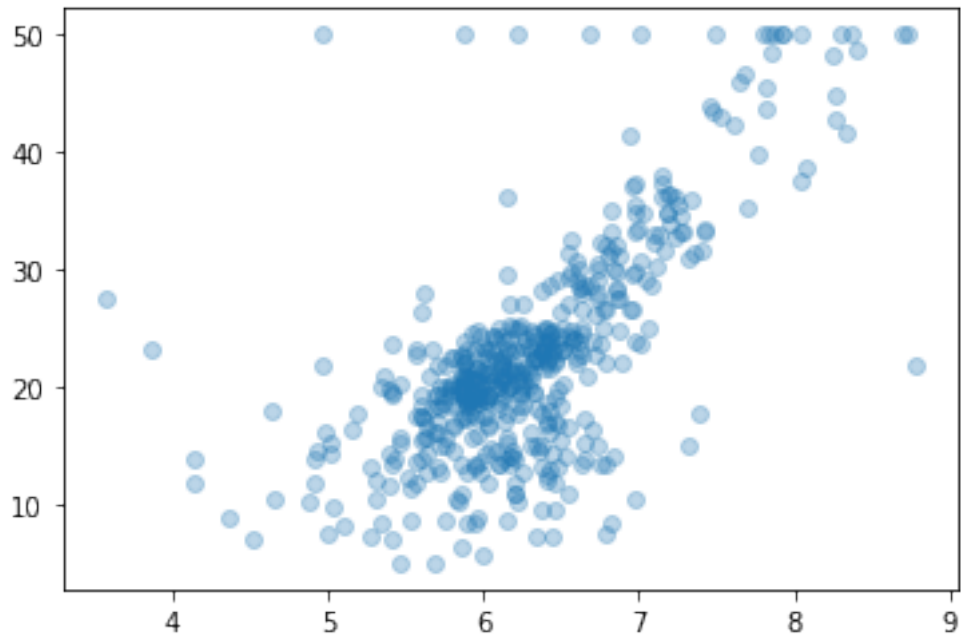
In [5]: boston = pd.DataFrame(boston_dataset.data, columns=boston_dataset.feature_names)
        boston['MEDV'] = boston_dataset.target
        boston.head()
```

Out [5]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	

	PTRATIO	B	LSTAT	MEDV
0	15.3	396.90	4.98	24.0
1	17.8	396.90	9.14	21.6
2	17.8	392.83	4.03	34.7
3	18.7	394.63	2.94	33.4
4	18.7	396.90	5.33	36.2

```
In [6]: scatter(boston['RM'], boston['MEDV'], alpha=0.3)
        show()
```



```
In [7]: X = array(boston['RM'])
        Y = array(boston['MEDV'])
```

```
In [8]: import torch
```

```
In [9]: print(X.shape, Y.shape)
```

```
(506,) (506,)
```

```
In [10]: X = torch.Tensor(X.reshape(506, 1))
          Y = torch.Tensor(Y.reshape(506, 1))
```

```
In [11]: print(X.shape, Y.shape)
```

```
torch.Size([506, 1]) torch.Size([506, 1])
```

```
In [12]: class LinearRegretion(torch.nn.Module):
        def __init__(self, input_dim):
            super().__init__()

            self.F = torch.nn.Linear(input_dim, 1)
            self.loss = None
            self.accuracy = None

        def forward(self, x):
```

```

        return self.F(x)

    def fit(self, x, y, epochs=1, lr=0.01):
        loss_fn = torch.nn.MSELoss()
        optimizer = torch.optim.SGD(self.parameters(), lr=lr)

        self.train()

        for i in range(0, epochs):
            y_ = self.forward(x)
            loss = loss_fn(y_, y)


            optimizer.zero_grad()
            loss.backward()
            optimizer.step()

        self.loss = loss.detach().numpy()
        self.accuracy = self.loss / x.shape[0]

```

```
In [13]: h = LinearRegretion(1)
```

```
In [14]: hy = h(X).detach()
scatter(X, Y, alpha=0.3)
plot(X, hy, c="brown")
show()
```



```
In [15]: print(h.F.weight)
```


Parameter containing:
tensor([[0.4203]], requires_grad=True)

```
In [16]: print(h.F.bias)
```

Parameter containing:
tensor([0.2636], requires_grad=True)

```
In [17]: h.fit(X, Y, epochs=100000, lr=0.003)
```

```
hy = h(X).detach()
scatter(X, Y, alpha=0.3)
plot(X, hy, c="brown")
show()
```



```
In [18]: for param in h.parameters():
          print(param)
```

Parameter containing:
tensor([[9.0979]], requires_grad=True)
Parameter containing:
tensor([-34.6439], requires_grad=True)

```
In [19]: print(h.F.weight)
```

```
Parameter containing:
tensor([[9.0979]], requires_grad=True)
```

```
In [20]: print(h.F.bias)
```

```
Parameter containing:
tensor([-34.6439], requires_grad=True)
```

```
In [21]: print("Loss: ", h.loss)
          print("Accuracy: ", h.accuracy)
```

```
Loss:  43.60056
Accuracy:  0.08616711311189554
```

```
In [22]: print((h(torch.Tensor([6]))*1000).detach().numpy())

[19943.55]
```

```
In [23]: print(h)
```

```
LinearRegretion(
  (F): Linear(in_features=1, out_features=1, bias=True)
)
```


```
In [24]: print(h.state_dict())

OrderedDict([('F.weight', tensor([[9.0979]])), ('F.bias', tensor([-34.6439]))])
```

```
In [25]: torch.save(h.state_dict(), "linear_state_dict.pt")
```

```
In [26]: other_model = LinearRegretion(1)
```

```
In [27]: hy = other_model(X).detach()
          scatter(X, Y, alpha=0.3)
          plot(X, hy, c="brown")
          show()
```



```
In [28]: other_model.load_state_dict(torch.load("linear_state_dict.pt"))
```

```
Out[28]: <All keys matched successfully>
```

```
In [29]: hy = other_model(X).detach()
scatter(X, Y, alpha=0.3)
plot(X, hy, c="brown")
show()
```

